



# COMP3311 Week 2 Wednesday Lecture

---

- Week 02 Wednesday
- Assignment 1 Database
- Recap
- Exercise: ER-to-SQL for Beer Database
- Mapping Composite Attributes
- Mapping Multi-valued Attributes (MVAs)
- Mapping Subclasses
- Exercise: ER-to-SQL (1)
- Exercise: ER-to-SQL (2)
- Exercise: ER-to-SQL (3)
- PostgreSQL Databases
- **psql**
- Exercise: Creating a database
- Populating a Database
- Exercise: Inserting Tuples
- Exercise: More Inserting Tuples
- Bulk Insertion
- Dump/Restore
- Exercise: Playing with Beer Database

## ❖ Week 02 Wednesday

---

### In today's lecture ...

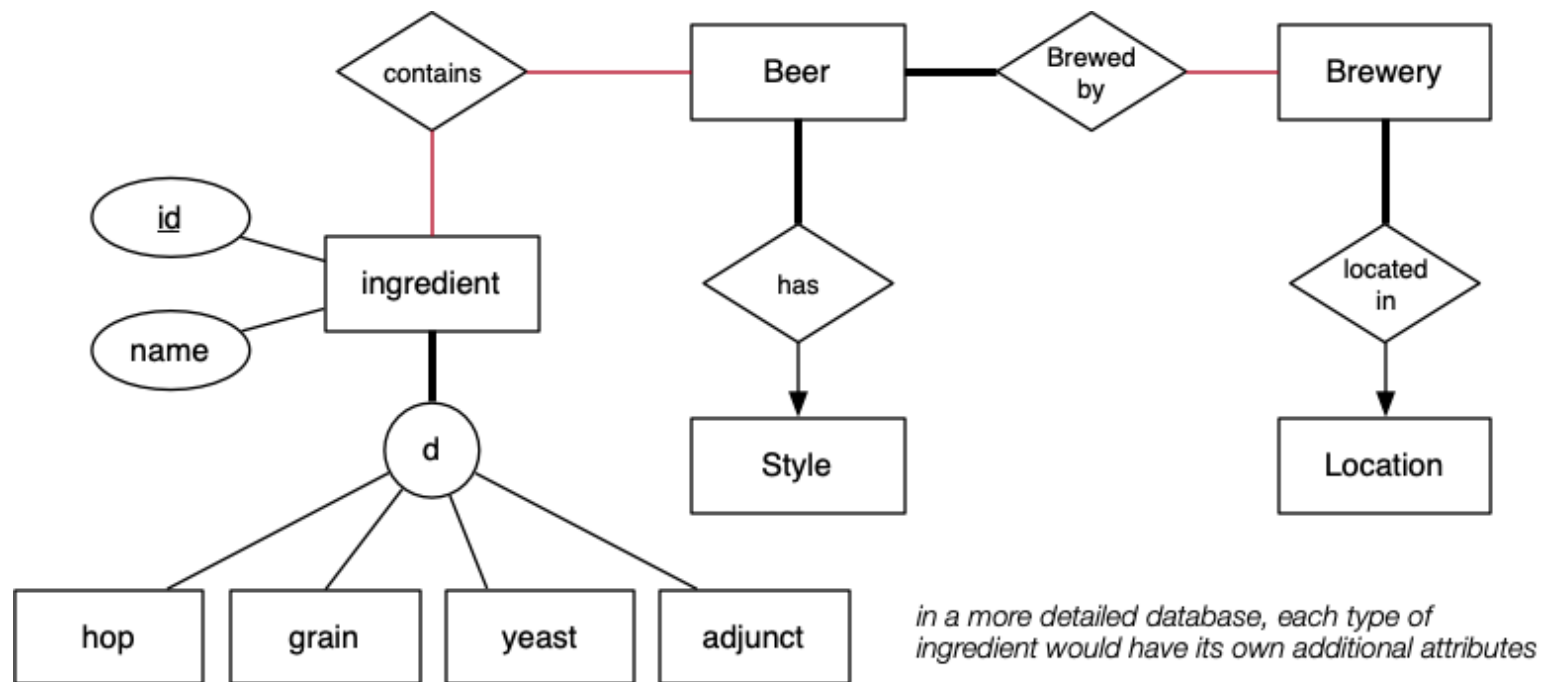
- More ER → SQL, and SQL DDL
- Building/restoring a database

### Things to do ...

- Quiz before Friday midnight
- Set up your PostgreSQL server  
(300 students have logged in to vxdb2 and have /localstorage)
- Help Session, Friday 4pm, Location: TBA

## ❖ Assignment 1 Database

Database about beer and breweries



## ❖ Assignment 1 Database (cont)

Details of entities, with example data (does not include all entities):

```
Beers(id, name, brewed, style, abv, ibu, sold_in, volume, notes, rating)
```

```
(123, 'VB', 2020, *Lager, 5.0, 30, can, 375, 'Worst beer in world', 1)
```

```
Brewers(id, name, founded, website, located_in)
```

```
(321, 'Carlton', 1899, 'www.carlton.com.au', *Melbourne)
```

```
Styles(id, name, min_abv, max_abv)
```

```
(456, 'Lager', 4.0, 6.0)
```

```
Ingredients(id, itype, name)
```

```
(654, 'hop', 'Cascade')
```

```
Brewed_by(beer, brewery)
```

```
(*VB, *Carlton) ... represented as (123, 321)
```



## ❖ Recap

---

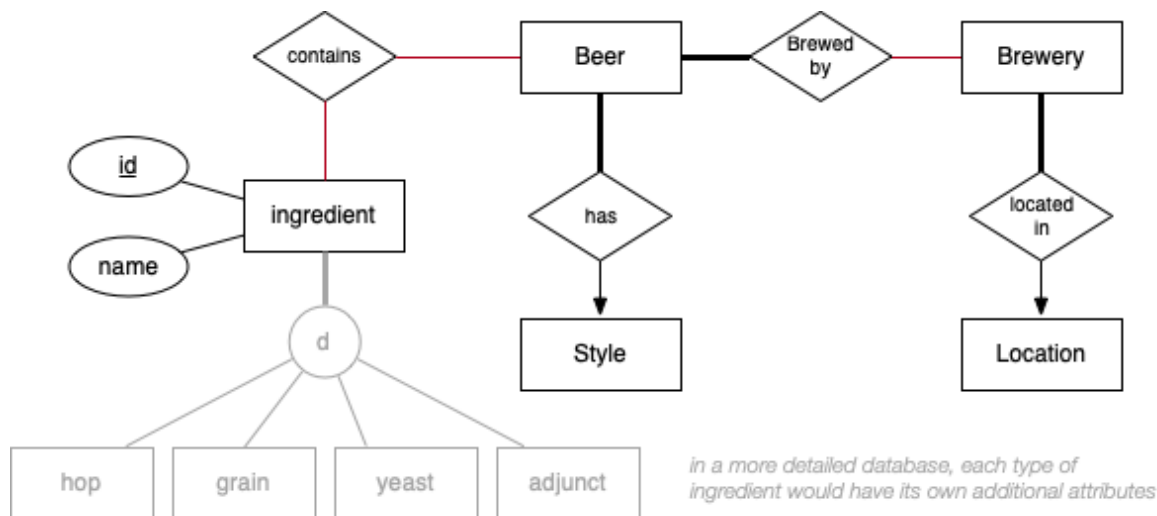
### ER → Relational/SQL Mapping

- attributes → attributes
- entities → relations/tables
- 1:1 relationships → foreign key
- 1:n relationships → foreign key
- n:m relationships → link table
- composite attributes → attributes
- multi-valued attributes → table

## ❖ Exercise: ER-to-SQL for Beer Database

Convert the beer ER data model to an SQL schema


- assume that each entity has attributes **id** and **name**
- treat **Ingredient** as simple entity; ignore sub-classes



## ❖ Mapping Composite Attributes

Composite attributes are mapped by concatenation or flattening.

Example:



### *Relational Version #1*

Person

<b>SSN</b>	name	address
------------	------	---------

### *Relational Version #2*

Person


SSN	given	family	.....	
.....	number	street	suburb	pcode



## ❖ Mapping Multi-valued Attributes (MVAs)

MVAs are mapped by a new table linking values to their entity.

Example:



### *Relational Version*

<i>Person</i>	<b>SSN</b>	name	birthday
---------------	------------	------	----------

<i>FavColour</i>	<b>SSN</b>	colour
------------------	------------	--------

## ❖ Mapping Subclasses

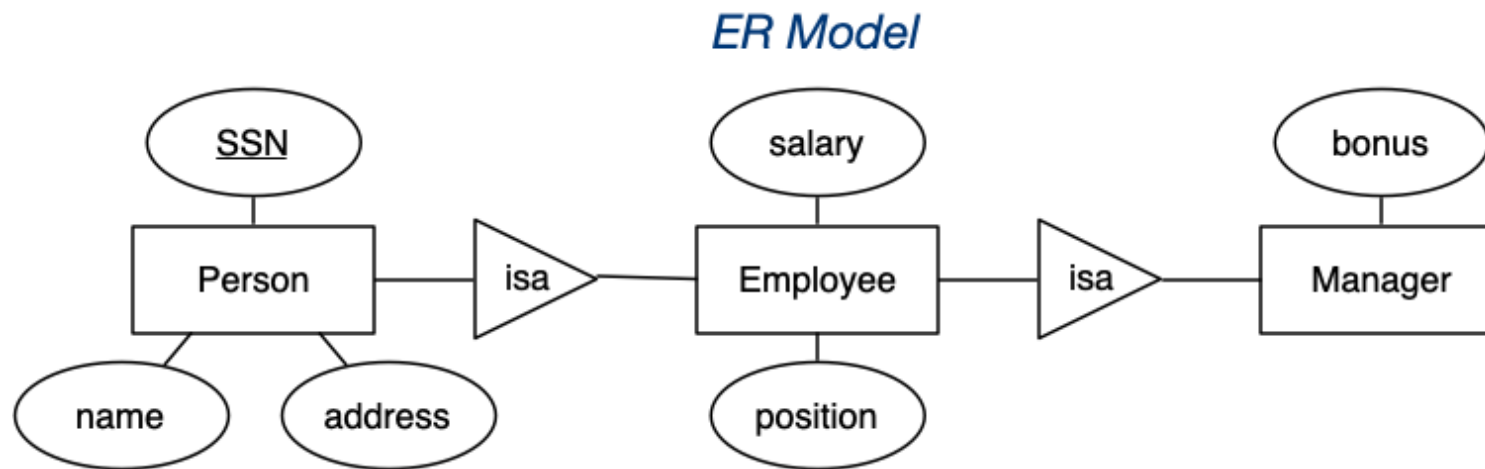
Three different approaches to mapping subclasses to tables:

- ER style
  - each entity becomes a separate table,
  - containing attributes of subclass + FK to superclass table
- object-oriented
  - each entity becomes a separate table,
  - inheriting all attributes from all superclasses
- single table with nulls
  - whole class hierarchy becomes one table,
  - containing all attributes of all subclasses (null, if unused)

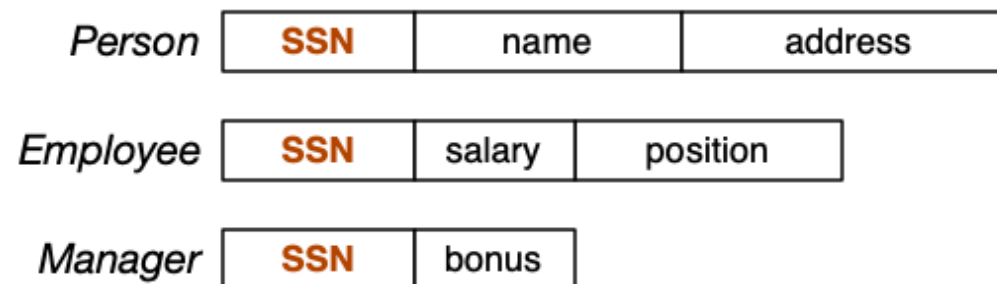
Which mapping is best depends on how data is to be used.

## ❖ Mapping Subclasses (cont)

Example of ER-style mapping:



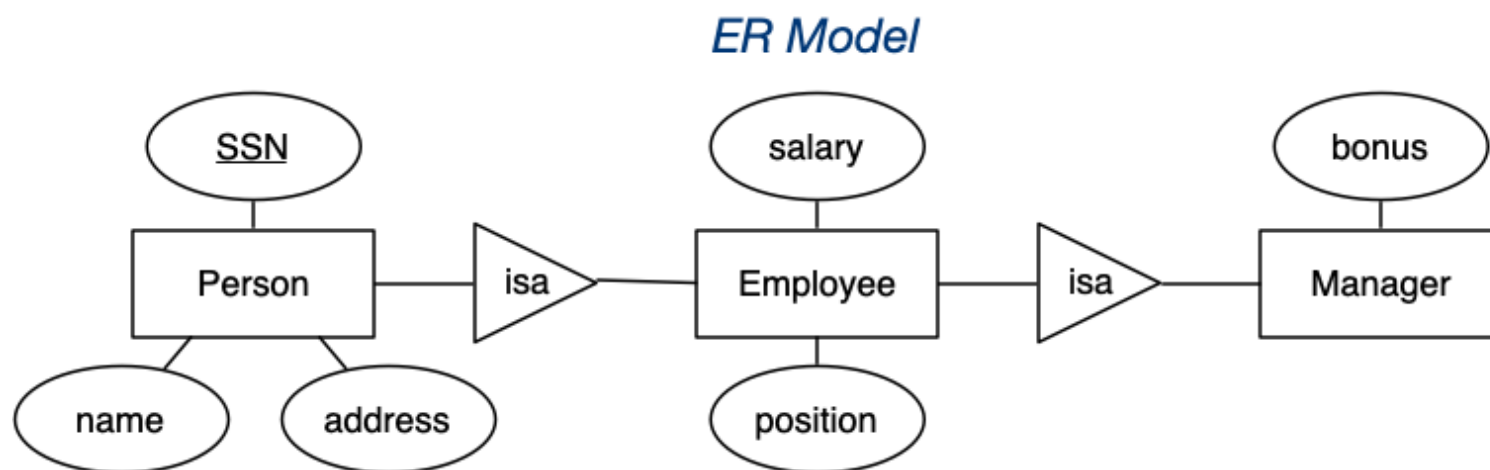
### *Relational Version*





## ❖ Mapping Subclasses (cont)

Example of object-oriented mapping:



### *Relational Version*

<i>Person</i>	<b>SSN</b>	name	address
---------------	------------	------	---------

<i>Employee</i>	<b>SSN</b>	name	address	salary	position
-----------------	------------	------	---------	--------	----------


<i>Manager</i>	<b>SSN</b>	name	address	salary	position	bonus
----------------	------------	------	---------	--------	----------	-------




## ❖ Mapping Subclasses (cont)

Example of single-table-with-nulls mapping:

*ER Model*



*Relational Version*

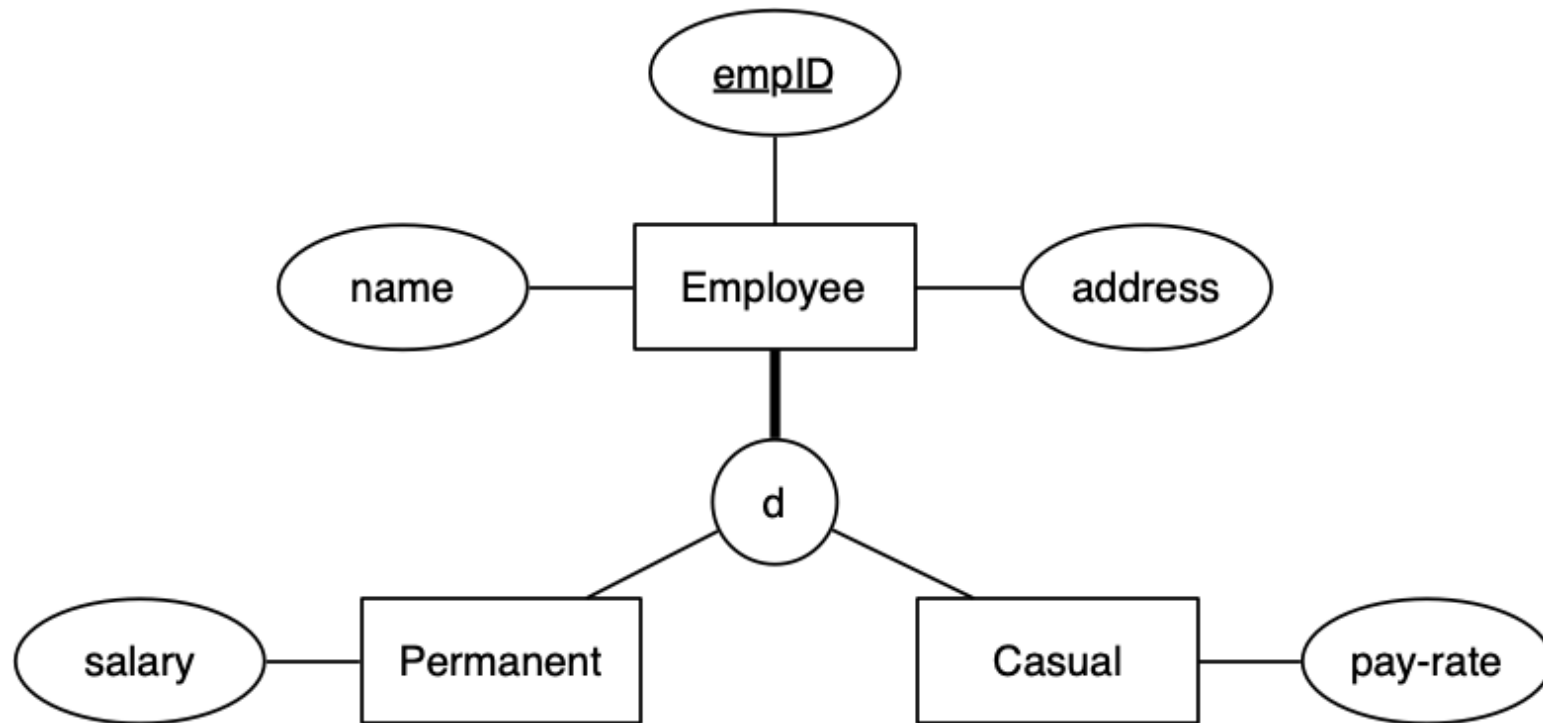


*NULL for Person who is not Employee*

*NULL for Employee who is not Manager*

## ❖ Exercise: ER-to-SQL (1)

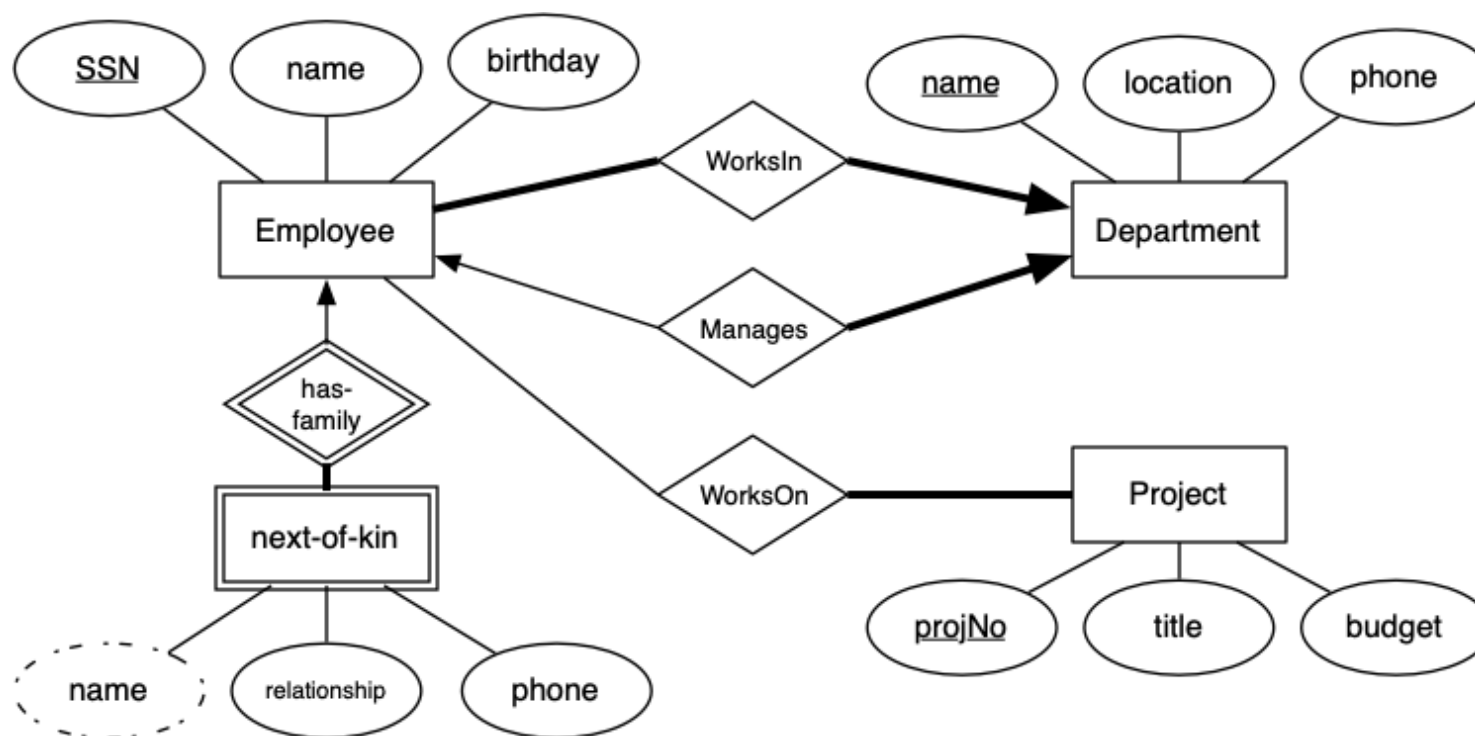
Convert the following class hierarchy to SQL using ER mapping:





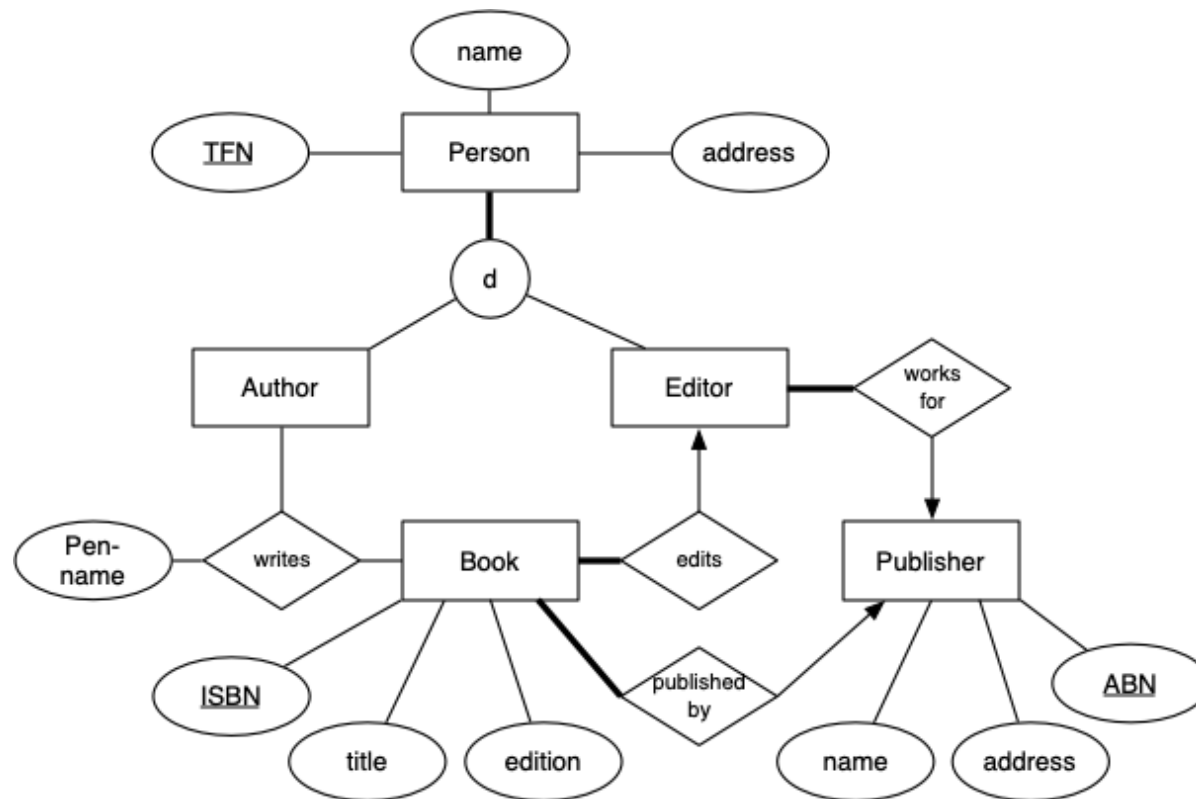
## ❖ Exercise: ER-to-SQL (2)

Convert this ER design to SQL:



## ❖ Exercise: ER-to-SQL (3)

Convert the Publishing ER model to SQL



- A TFN is stored as a 9-digit number
- An ABN is stored as an 11-digit number

- An ISBN (13-digit version) looks like 978-3-16-148410-0

## ❖ PostgreSQL Databases

Create a database in PostgreSQL via

```
$ createdb DatabaseName
```

Creates an empty database (no schema, no data)

Remove a database in PostgreSQL via

```
$ dropdb DatabaseName
```

Removes schema and all data permanently

Remove an entire PostgreSQL server (on **vxdb2**)

```
$ rm -fr /localstorage/$USER/pgsql
```

Removes all server files, all databases, all data !





The **psql** command is a shell that allows you to

- connect to databases (one at a time)
- ask SQL queries on a database
- find information (meta data) about a database
- add/delete/update tuples in tables

Usage:

```
$ psql mydb
...
mydb=# \d
...
mydb=$ select * from SomeTable;
...
mydb=# \q
```



## ❖ **psql** (cont)

A useful way to use **psql**:

```
$ psql -l
```

Gives a list of all databases under your PostgreSQL server.

The "databases" **postgres**, **template1**, **template2**

- are special databases used internally by PostgreSQL
- do not **dropdb** them





## ❖ **psql** (cont)

---

The **psql** command has several prompts

- **db=#** ... waiting to start a command
- **db-#** ... waiting for rest of command
- **db(#** ... waiting for rest of expression ( . . . )
- **db' #** ... waiting to finish a string (expecting closing ' )

Note that **db** will be (replaced by) the name of the current database

Note that **#** means you are super-user; normal users get **>**

SQL statements can span several lines, terminated by typing **;**

**psql** meta-commands are single-line commands

## ❖ **psql** (cont)

---

**psql** has a range of **meta-commands**, beginning with **\**

- **help** ... a quick list of useful meta-commands
- **\?** ... a list of all meta-commands (very many of them)
- **\d** ... list of all tables and views in current schema
- **\dt** ... list of all tables in current schema
- **\d Table** ... list of all attributes in *Table*
- **\df** ... list of all user-defined functions in current schema
- **\q** ... exit **psql**

## ❖ Exercise: Creating a database

---

On **vxdb2**, do the following:

- create a database called **xyz**
- examine its (empty) schema
- within **xyz** create a table

```
create table R (  
    x integer primary key,  
    y float not null,  
    z text  
);
```

- examine the schema again
- examine the attributes of table **R**
- how many tuples are in table **R**?
- remove the database **xyz**



## ❖ Populating a Database

Basic way of adding tuples to a database:

```
db=# insert into Table values (val1,val2,...);
```

Adds a tuple to table ***Table*** assuming

- values satisfy all constraints on tuples/table

Ways it can fail ...

- value for primary key field is already in the table
- tuple has **null** values for fields defined as **not null**
- value for some field violates constraints on that field
- etc. etc. etc.

## ❖ Exercise: Inserting Tuples

Which **insert** statements are successful?

If successful, what tuple value is inserted?

```
create type Mood as enum ('sad','happy');

create table People (
    name text not null,
    feels Mood
);

insert into People values ('John','happy');
insert into People values ('Andrew','angry');
insert into People values ('Tina',null);
insert into People(name) values ('Anne');
insert into People(feels) values ('happy');
```

## ❖ Exercise: More Inserting Tuples

Which **insert** statements are successful?

If successful, what tuple value is inserted?

```
create domain PosInt as integer check (value > 0);

create table Points (
    x PosInt default 1,
    y posint
);

insert into Points values (3,4);
insert into Points values (3,null);
insert into Points values (-3,4);
insert into Points(y) values (5);
insert into Points values (default,5);
```

## ❖ Bulk Insertion

Entering tuples interactively one-by-one is not feasible

Alternative: put **insert** statements in a file and run

```
$ psql DatabaseName -f FileName
```

Attempts to execute each **insert** statement:

- if tuple valid, inserted into database
- if tuple not valid, prints error message, then continues

Note that ***FileName*** can contain any SQL statements

Might consist only of **create table** statements to build a schema



## ❖ Bulk Insertion (cont)

A common way of building a database

```
$ createdb mydb  
$ psql mydb -f schema.sql  
$ psql mydb -f data.sql
```

where

- **schema.sql** contains table and type definitions
- **data.sql** contains statements to insert tuples

## ❖ Bulk Insertion (cont)

How I "debug" a database schema

```
$ dropdb mydb
$ createdb mydb
$ psql mydb -f schema.sql > .errs 2>&1
$ vi .errs
# fix any errors that appear in .errs
$ psql mydb -f data.sql > .errs 2>&1
$ vi .errs
# fix any errors that appear in .errs
```

Repeat until **.errs** contains no lines with **ERROR**

## ❖ Bulk Insertion (cont)

Alternative way of inserting tuples

```
copy TableName ( attribute names ) from stdin;  
... lines containing tab-separated values ...  
... one value for each of the named attributes ...  
\.
```

Difference between **copy** and multiple **inserts**

- with **insert** ...
  - all tuples with valid values are inserted
  - tuples with invalid values are not inserted
- with **copy** ...
  - if any tuple contains invalid values, nothing is inserted

## ❖ Dump/Restore

Once a database is built, can make a complete copy in a text file

- the whole schema (including types, constraints, etc), and all data

by running the command

```
$ pg_dump -0 -x DatabaseName > DumpFileName
```

and can make a new copy via

```
$ createdb newdb  
$ psql newdb -f DumpFileName
```

We generally supply databases using pre-built dump files

## ❖ Exercise: Playing with Beer Database

---

Load up the database from **ass1.dump**

Guess some SQL to answer the following:

- what tables are in the database
- what fields/attributes are in the **Beers** table
- what fields/attributes are in the **Breweries** table
- how many beers there are
- how many breweries there are
- what beers have "Black Lung" in their name

---

Produced: 20 Sep 2023